

Chandan Mohan

linkedin.com/in/imchandanmohan

Email : chandan.mohan@gwu.edu

Mobile : +1-240-990-9298

EDUCATION

- **George Washington University** Washington, D.C.
Master of Science in Computer Science; GPA: 3.80 Dec. 2023 – May 2024
Graduate Teaching Assistant (CSCI 6442 - Advanced DBMS II): Worked on Terraform scripts in AWS to automate infrastructure deployment for distributed systems and data pipelines, enhancing expertise in data warehouse design.

EXPERIENCE

- **Advaiya Solutions** Udaipur, IN
Consultant Jan 2023 – Oct 2023
 - **TensorFlow**: Worked on implementing deep learning models using TensorFlow for structured and unstructured data. Focused on model optimization and deployment.
 - **Apache Beam**: Built scalable data pipelines using Apache Beam for real-time and batch processing across Google Cloud Platform.
- **VNB Consulting Services** Bangalore, IN
Senior Software Engineer Jun 2020 – Dec 2022
 - **Notifications**: Built services for email, push, and in-app notifications with features like delivery time optimization, tracking, queuing, and A/B testing.
 - **Nostos**: Developed a bulk data processing and injection system from Hadoop to Cassandra with a REST API layer for serving offline data.
 - **Workflows**: Used Dataduct (open-source) for building reusable and automated data pipelines, improving engineering efficiency.
 - **Data Collection**: Designed an internal crowdsourcing platform to embed and manage surveys across the Coursera ecosystem.
 - **Dev Environment**: Set up Docker-based analytics environments on AWS, standardizing Python and R libraries used by data science teams.
 - **Recommendations**: Contributed to core recommendation systems for Coursera's homepage and discovery. Supported both training and real-time inference.
 - **Content Discovery**: Improved search and browsing experience with onboarding personalization, ranking improvements, and indexing optimization.

PROJECTS

- **Credit Risk Prediction System**:
 - Developed an ML pipeline using Logistic Regression and XGBoost to predict credit risk with 87% accuracy. Performed feature engineering and hyperparameter tuning, improving AUC-ROC by 12%. Contributed to core recommendation systems for Coursera's homepage and discovery. Supported both training and real-time inference.
 - Deployed model via Flask API and integrated into a dashboard for real-time decision-making. Applied augmentation techniques to reduce overfitting and cut training time by 86%. Built services for email, push, and in-app notifications. Designed an internal crowdsourcing platform to embed and manage surveys across the Coursera ecosystem.
- **Image Classification using Transfer Learning**:
 - Fine-tuned ResNet50 and EfficientNet models to classify medical images with 92% validation accuracy. Set up Docker-based analytics environments on AWS, standardizing Python and R libraries used by data science teams. Built services for email, push, and in-app notifications with features like delivery time optimization, tracking, queuing, and A/B testing.
 - Used MLflow for experiment tracking and versioning to ensure reproducibility. Applied augmentation techniques to reduce overfitting and cut training time by 50%.

SKILLS

- **Machine Learning**: TensorFlow, PyTorch, scikit-learn, XGBoost, LightGBM
- **Data Engineering**: SQL, Apache Spark, Pandas, Airflow
- **Cloud & Deployment**: AWS, Docker, Kubernetes, Flask
- **Languages**: Python, Java, C++