# Chandan Mohan

linkedin.com/in/imchandanmohan

Email : chandan.mohan@gwu.edu
Mobile : +1-240-990-9298


## Education

- **George Washington University** — Washington, D.C.
  *Master of Science in Computer Science;* ***GPA: 3.80*** — *Dec. 2023 – May 2024*
  **Graduate Teaching Assistant** (CSCI 6442 - Advanced DBMS II): Worked on Terraform scripts in AWS to automate infrastructure deployment for distributed systems and data pipelines, enhancing expertise in data warehouse design.

## Experience

- **Advaiya Solutions** — Udaipur, IN
  *Consultant* — *Jan 2023 – Oct 2023*
  - **TensorFlow**: Worked on implementing deep learning models using TensorFlow for structured and unstructured data. Focused on model optimization and deployment.
  - **Apache Beam**: Built scalable data pipelines using Apache Beam for real-time and batch processing across Google Cloud Platform.
- **VNB Consulting Services** — Bangalore, IN
  *Senior Software Engineer* — *Jun 2020 – Dec 2022*
  - **Notifications**: Built services for email, push, and in-app notifications with features like delivery time optimization, tracking, queuing, and A/B testing.
  - **Nostos**: Developed a bulk data processing and injection system from Hadoop to Cassandra with a REST API layer for serving offline data.
  - **Workflows**: Used Dataduct (open-source) for building reusable and automated data pipelines, improving engineering efficiency.
  - **Data Collection**: Designed an internal crowdsourcing platform to embed and manage surveys across the Coursera ecosystem.
  - **Dev Environment**: Set up Docker-based analytics environments on AWS, standardizing Python and R libraries used by data science teams.
  - **Recommendations**: Contributed to core recommendation systems for Coursera's homepage and discovery. Supported both training and real-time inference.
  - **Content Discovery**: Improved search and browsing experience with onboarding personalization, ranking improvements, and indexing optimization.

## Projects

- **Credit Risk Prediction System**:
  - Developed an ML pipeline using Logistic Regression and XGBoost to predict credit risk with 87% accuracy. Performed feature engineering and hyperparameter tuning, improving AUC-ROC by 12%. Contributed to core recommendation systems for Coursera's homepage and discovery. Supported both training and real-time inference.
  - Deployed model via Flask API and integrated into a dashboard for real-time decision-making. Applied augmentation techniques to reduce overfitting and cut training time by 86%. Built services for email, push, and in-app notifications. Designed an internal crowdsourcing platform to embed and manage surveys across the Coursera ecosystem.
- **Image Classification using Transfer Learning**:
  - Fine-tuned ResNet50 and EfficientNet models to classify medical images with 92% validation accuracy.Set up Docker-based analytics environments on AWS, standardizing Python and R libraries used by data science teams.Built services for email, push, and in-app notifications with features like delivery time optimization, tracking, queuing, and A/B testing.
  - Used MLflow for experiment tracking and versioning to ensure reproducibility. Applied augmentation techniques to reduce overfitting and cut training time by 50%.

## Skills

- **Machine Learning:** TensorFlow, PyTorch, scikit-learn, XGBoost, LightGBM
- **Data Engineering:** SQL, Apache Spark, Pandas, Airflow
- **Cloud & Deployment:** AWS, Docker, Kubernetes, Flask
- **Languages:** Python, Java, C++